

Mumble protocol 1.2.X reference (WIP)

Stefan Hacker

October 15, 2010

DISCLAIMER

THIS DOCUMENTATION IS PROVIDED BY THE MUMBLE PROJECT "AS IS" AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL THE MUMBLE PROJECT BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS DOCUMENTATION, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.

Contents

1	Introduction	3
2	Overview	3
3	Protocol stack (TCP)	4
4	Establishing a connection	5
4.1	Connect	5
4.2	Version exchange	6
4.3	Crypt setup	6
4.4	Authenticate	6
4.5	Channel states	7
4.6	User states	7
4.7	Server sync	7
4.8	Ping	8
5	Voice data	8
5.1	Enabling the UDP channel	8
5.2	Data	8
5.3	TCP tunnel	10
5.4	Encryption	10
5.5	PacketDataStream	10
6	This document is WIP	10
A	Appendix	11

A.1	Mumble.proto	11
-----	------------------------	----

1 Introduction

This document is meant to be a reference for the Mumble VoIP 1.2.X server-client communication protocol. It reflects the state of the protocol implemented in the Mumble 1.2.2 client and might be outdated by the time you are reading this. Be sure to check for newer revisions of this document on our website <http://www.mumble.info> . At the moment this document is work in progress.

2 Overview

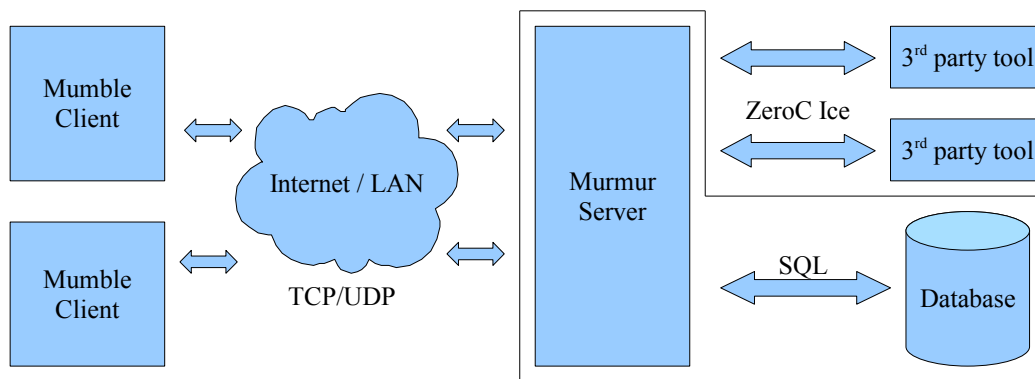


Figure 1: Mumble system overview

Mumble is based on a standard server-client communication model. It utilizes two channels of communication, the first one is a TCP connection which is used to reliably transfer control data between the client and the server. The second one is a UDP connection which is used for unreliable, low latency transfer of voice data.

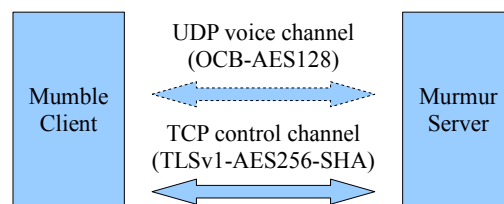


Figure 2: Mumble crypto types

Both are protected by strong cryptography, this encryption is mandatory and cannot be disabled. The TCP control channel uses TLSv1 AES256-SHA¹ while the voice channel

¹http://en.wikipedia.org/wiki/Transport_Layer_Security

is encrypted with OCB-AES128².

While the TCP connection is mandatory the UDP connection can be compensated by tunnelling the UDP packets through the TCP connection as described in the protocol description later.

3 Protocol stack (TCP)

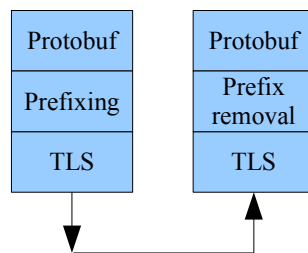


Figure 3: Mumble protocol stack

Mumble has a shallow and easy to understand stack. Basically it uses Googles Protocol Buffers³ with simple prefixing to distinguish the different kinds of packets sent through an TLSv1 encrypted connection. This makes the protocol very easily expandable.




Figure 4: Mumble packet

The prefix consists out of the two bytes defining the type of the packet in the payload and 4 bytes stating the length of the payload in bytes followed by the payload itself. The following packet types are available in the current protocol and all but UDPTunnel are simple protobuf messages. If not mentioned otherwise all fields are little-endian encoded.

For raw representation of each packet type see the attached Mumble.proto file.

²<http://www.cs.ucdavis.edu/~rogaway/ocb/ocb-back.htm>

³<http://code.google.com/p/protobuf/>

4 Establishing a connection

The following section is going to describe the communication between the server and the client during connection establishing, note that the first part of this section only contains the procedures for the TCP connection. After this the client will be visible to the other clients on the server and able to send other types of messages.

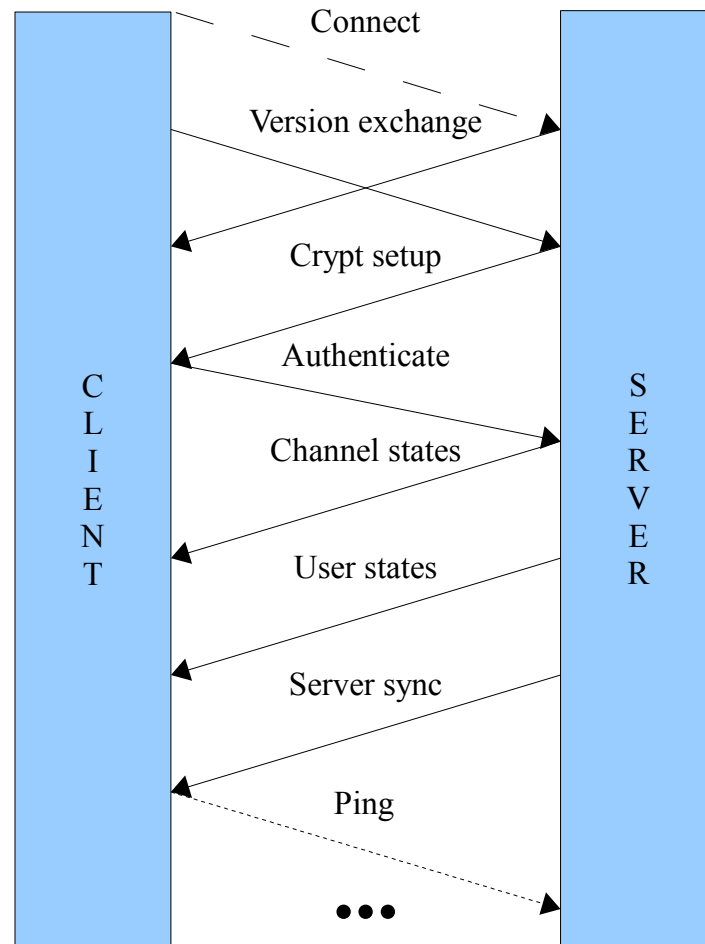


Figure 5: Mumble connection setup

4.1 Connect

As a basis for the synchronization procedure you first have to establish the TCP connection to the server and do a common TLSv1 handshake. To be able to use the complete feature set of the Mumble protocol it is recommended that your client provides a

strong certificate to the server. This however is not mandatory, you can connect to the server without providing a certificate, we do recommend to check the servers certificate though.

4.2 Version exchange

Once the TLS handshake is completed the server will send a Version packet to the client containing following information:

Version	
version	uint32
release	string
os	string
os_version	string

The client is supposed to send a Version packet with his information before any other messages. The version field of the packet contains the (major, minor, patch) tuple (e.g. 1.2.0) encoded like described in the following figure.

Major	Minor	Patch
2B	1B	1B

The release, os and os_version fields are common strings containing additional information about the client. This information is not interpreted in any way at the moment.

4.3 Crypt setup

Once the Version packets are exchanged the server will send a CryptSetup packet to the client. It contains the necessary cryptographic information to establish the OCB-AES128 encrypted UDP Voice channel. This will be described later in the section.

CryptSetup	
Key	ByteString
ServerNonce	ByteString
ClientNonce	ByteString

4.4 Authenticate

Before the client can be synchronized with the server state it has to authenticate itself to the server. This is done by sending an Authenticate packet.

Authenticate	
username	string
password	string
tokens	repeated string

The username and password are encoded as simple strings. Be aware that the server can impose restrictions on the username, also once the client registered a certificate with the server this field is only displayed in brackets behind the name the client possessed when he registered, for more information see the server documentation. The password must only be provided if the server is passworded, the client provided no certificate but wants to authenticate to an account which has a password set, or to access the SuperUser account. The third field called tokens contains a list of zero or more strings called tokens which are basically password which can give you access to a certain ACL group without actually being a registered member in them, again see the server documentation for more information.

4.5 Channel states

After the client is successfully authenticated the server starts synchronizing the state by transmitting a ChannelState message for every channel on this server. Note that these do not yet contain channel links. These are transmitted as updated directly after every channel has been transmitted. It contains the following information:

For more information please refer to Mumble.proto in the appendix.

4.6 User states

When the channels are synchronized the server send a UserState message for every user connected to the client containing the following data:

For more information please refer to Mumble.proto in the appendix.

4.7 Server sync

The client has now received a copy of the parts of the server state he needs to know about. To complete the synchronization the server transmits a ServerSync message containing the session id of the clients session, the maximum bandwidth allowed on this server, the servers welcome text as well as the permissions the client has in the channel he ended up.

For more information please refer to Mumble.proto in the appendix.

4.8 Ping

If the client wishes to maintain the connection to the server it is required to ping the server. If the server does not receive a ping for TODO seconds it will disconnect the client.

5 Voice data

5.1 Enabling the UDP channel

Before the UDP channel can reliably be used both sides should be certain that the connection works. Before the server may use the UDP connection to the client the client must first open a UDP socket and communicate its address to the server by sending a packet over UDP. Once the server has received an UDP transmission the server should start using the UDP channel for the voice packets. Respectively the client should not use the UDP channel for voice data until it is certain that the packets go through to the server.

In practice these requirements are filled with UDP ping. When the server receives a UDP ping packet (See figure 6) from the client it echoes the packet back. When the client receives this packet it can ascertain that the UDP channel works for two-way communication.

byte	:	type/flags	0010 0000 for Ping
varint	:	timestamp	Timestamp for the client.

Figure 6: UDP Ping packet

If the client stops receiving replies to the UDP packets at some point or never receives the first one it should immediately start tunneling the voice communication through TCP as described in section 5.3. When the server receives a tunneled packet over the TCP connection it must also stop using the UDP for communication. The client may continue sending UDP ping packets over the UDP channel and the server must echo these if it receives them. If the client later receives these echoes it may switch back to the UDP channel for voice communication. When the server receives a UDP voice communication packet from the client it should stop tunneling the packets as well.

5.2 Data

The voice data is transmitted in variable length packets that consist of header portion, followed by repeated data segments and an optional position part. The full packet

Header	byte : type/target	Bit 1-3: Type, Bit 4-8: Target
	varint : session	The session number of the source user
	varint : sequence	
Audio Repeated	byte : header	Bit 1: Terminator, Bit 2-8: Data length
	byte[] : data	Encoded voice frames
Position Optional	float : Pos 1	Positional audio positions
	float : Pos 2	Uses PacketDataStream encoding
	float : Pos 3	

Figure 7: UDP Voice packet

structure is shown in figure 7. The protocol transfers 64-bit integers using variable length encoding. This encoding is specified in section ??.

The first byte of the header contains the packet type and additional target specifier. The type is stored in the first three bits and specifies the type and encoding of the packet. Current types are listed in table ?. The remaining 5 bits specify additional packet-wide options. For voice packets the values specify the voice target as listed in table 2.

Table 1: UDP Types

Type	Description
0	CELT Alpha encoded voice data
1	Ping packet (See section 5.1)
2	Speex encoded voice data
3	CELT Beta encoded voice data
4-8	Unused

Table 2: UDP targets

Target	Description
0	Normal talking
1	Whisper to channel
2-30	Direct whisper (Refer to VoiceTarget, ??). Always 2 for incoming whisper.
31	Server loopback

The audio frames consist of one byte long header and up to 127 bytes long data portion. The first bit in the header is the **Terminator bit** which informs the receiver whether there are more audio frames after this one. This bit is turned on (value 1) for all but the last frame in the current UDP packet. Rest of the seven bits in the header specify the length of the data portion. The data portion is encoded using one of the supported

codecs. The exact codec is specified in the type portion of the whole packet (See table 1). *The data in each frame is encoded separately.*

5.3 TCP tunnel

When the UDP packets are tunneled through the TCP tunnel they are prefixed with the TCP protocol header that contains the packet type and length and sent through the connection. (Figure 8)

Type	Length	UDP Packet
1B	3B	0-2048 KiB

Figure 8: UDP Voice packet

5.4 Encryption

All the voice packets are encrypted once during transfer. The actual encryption depends on the used transport layer. If the packets are tunneled through TCP they are encrypted using the TLS that encrypts the whole TCP connection and if they are sent directly using UDP they must be encrypted using the OCB-AES128 encryption. The OCB-AES128 encryption is described in section ??.

When implementing the protocol it is easier to ignore the UDP transfer layer at first and just tunnel the UDP data through the TCP tunnel. The TCP layer must be implemented for authentication in any case. Making sure that the voice transmission works before implementing the UDP protocol simplifies debugging greatly. The UDP protocol is a required part of the specification though.

5.5 PacketDataStream

The PacketDataStream class is used to serialize/deserialize the data packets received on the UDP connection or via the TCP-Tunneling. As the name implies it provides a stream based access to the data it contains. To pull data from it the user has to know what is located on the current position in the stream (e.g. a uint32, utf8 string and so on), the class itself is not aware of it's contents.

6 This document is WIP

SORRY BUT THIS DOCUMENT IS WORK IN PROGRESS. AT THE MOMENT IT LACKS A LOT OF IMPORTANT INFORMATION BUT WE HOPE TO BE ABLE TO FINISH THIS DOCUMENT SOMEDAY :-)

A Appendix

A.1 Mumble.proto

```
1 package MumbleProto;
2
3 option optimize_for = SPEED;
4
5 message Version {
6     optional uint32 version = 1;
7     optional string release = 2;
8     optional string os = 3;
9     optional string os_version = 4;
10 }
11
12 message UDPTunnel {
13     required bytes packet = 1;
14 }
15
16 message Authenticate {
17     optional string username = 1;
18     optional string password = 2;
19     repeated string tokens = 3;
20     repeated int32 celt_versions = 4;
21 }
22
23 message Ping {
24     optional uint64 timestamp = 1;
25     optional uint32 good = 2;
26     optional uint32 late = 3;
27     optional uint32 lost = 4;
28     optional uint32 resync = 5;
29     optional uint32 udp_packets = 6;
30     optional uint32 tcp_packets = 7;
31     optional float udp_ping_avg = 8;
32     optional float udp_ping_var = 9;
```

```

33         optional float tcp_ping_avg = 10;
34         optional float tcp_ping_var = 11;
35     }
36
37     message Reject {
38         enum RejectType {
39             None = 0;
40             WrongVersion = 1;
41             InvalidUsername = 2;
42             WrongUserPW = 3;
43             WrongServerPW = 4;
44             UsernameInUse = 5;
45             ServerFull = 6;
46             NoCertificate = 7;
47         }
48         optional RejectType type = 1;
49         optional string reason = 2;
50     }
51
52     message ServerConfig {
53         optional uint32 max_bandwidth = 1;
54         optional string welcome_text = 2;
55         optional bool allow_html = 3;
56         optional uint32 message_length = 4;
57         optional uint32 image_message_length = 5;
58     }
59
60     message ServerSync {
61         optional uint32 session = 1;
62         optional uint32 max_bandwidth = 2;
63         optional string welcome_text = 3;
64         optional uint64 permissions = 4;
65     }
66
67     message ChannelRemove {
68         required uint32 channel_id = 1;
69     }
70
71     message ChannelState {
72         optional uint32 channel_id = 1;
73         optional uint32 parent = 2;
74         optional string name = 3;
75         repeated uint32 links = 4;

```

```

76         optional string description = 5;
77         repeated uint32 links_add = 6;
78         repeated uint32 links_remove = 7;
79         optional bool temporary = 8 [default = false];
80         optional int32 position = 9 [default = 0];
81         optional bytes description_hash = 10;
82     }
83
84     message UserRemove {
85         required uint32 session = 1;
86         optional uint32 actor = 2;
87         optional string reason = 3;
88         optional bool ban = 4;
89     }
90
91     message UserState {
92         optional uint32 session = 1;
93         optional uint32 actor = 2;
94         optional string name = 3;
95         optional uint32 user_id = 4;
96         optional uint32 channel_id = 5;
97         optional bool mute = 6;
98         optional bool deaf = 7;
99         optional bool suppress = 8;
100        optional bool self_mute = 9;
101        optional bool self_deaf = 10;
102        optional bytes texture = 11;
103        optional bytes plugin_context = 12;
104        optional string plugin_identity = 13;
105        optional string comment = 14;
106        optional string hash = 15;
107        optional bytes comment_hash = 16;
108        optional bytes texture_hash = 17;
109        optional bool priority_speaker = 18;
110        optional bool recording = 19;
111    }
112
113    message BanList {
114        message BanEntry {
115            required bytes address = 1;
116            required uint32 mask = 2;
117            optional string name = 3;
118            optional string hash = 4;

```

```

119         optional string reason = 5;
120         optional string start = 6;
121         optional uint32 duration = 7;
122     }
123     repeated BanEntry bans = 1;
124     optional bool query = 2 [default = false];
125 }
126
127 message TextMessage {
128     optional uint32 actor = 1;
129     repeated uint32 session = 2;
130     repeated uint32 channel_id = 3;
131     repeated uint32 tree_id = 4;
132     required string message = 5;
133 }
134
135 message PermissionDenied {
136     enum DenyType {
137         Text = 0;
138         Permission = 1;
139         SuperUser = 2;
140         ChannelName = 3;
141         TextTooLong = 4;
142         H9K = 5;
143         TemporaryChannel = 6;
144         MissingCertificate = 7;
145         UserName = 8;
146         ChannelFull = 9;
147     }
148     optional uint32 permission = 1;
149     optional uint32 channel_id = 2;
150     optional uint32 session = 3;
151     optional string reason = 4;
152     optional DenyType type = 5;
153     optional string name = 6;
154 }
155
156 message ACL {
157     message ChanGroup {
158         required string name = 1;
159         optional bool inherited = 2 [default = true];
160         optional bool inherit = 3 [default = true];
161         optional bool inheritable = 4 [default = true];

```

```

162         repeated uint32 add = 5;
163         repeated uint32 remove = 6;
164         repeated uint32 inherited_members = 7;
165     }
166     message ChanACL {
167         optional bool apply_here = 1 [default = true];
168         optional bool apply_subs = 2 [default = true];
169         optional bool inherited = 3 [default = true];
170         optional uint32 user_id = 4;
171         optional string group = 5;
172         optional uint32 grant = 6;
173         optional uint32 deny = 7;
174     }
175     required uint32 channel_id = 1;
176     optional bool inherit_acls = 2 [default = true];
177     repeated ChanGroup groups = 3;
178     repeated ChanACL acls = 4;
179     optional bool query = 5 [default = false];
180 }
181
182 message QueryUsers {
183     repeated uint32 ids = 1;
184     repeated string names = 2;
185 }
186
187 message CryptSetup {
188     optional bytes key = 1;
189     optional bytes client_nonce = 2;
190     optional bytes server_nonce = 3;
191 }
192
193 message ContextActionAdd {
194     enum Context {
195         Server = 0x01;
196         Channel = 0x02;
197         User = 0x04;
198     }
199     required string action = 1;
200     required string text = 2;
201     optional uint32 context = 3;
202 }
203
204 message ContextAction {

```



```

205         optional uint32 session = 1;
206         optional uint32 channel_id = 2;
207         required string action = 3;
208     }
209
210     message UserList {
211         message User {
212             required uint32 user_id = 1;
213             optional string name = 2;
214         }
215         repeated User users = 1;
216     }
217
218     message VoiceTarget {
219         message Target {
220             repeated uint32 session = 1;
221             optional uint32 channel_id = 2;
222             optional string group = 3;
223             optional bool links = 4 [default = false];
224             optional bool children = 5 [default = false];
225         }
226         optional uint32 id = 1;
227         repeated Target targets = 2;
228     }
229
230     message PermissionQuery {
231         optional uint32 channel_id = 1;
232         optional uint32 permissions = 2;
233         optional bool flush = 3 [default = false];
234     }
235
236     message CodecVersion {
237         required int32 alpha = 1;
238         required int32 beta = 2;
239         required bool prefer_alpha = 3 [default = true];
240     }
241
242     message UserStats {
243         message Stats {
244             optional uint32 good = 1;
245             optional uint32 late = 2;
246             optional uint32 lost = 3;
247             optional uint32 resync = 4;

```

```

248     }
249
250     optional uint32 session = 1;
251     optional bool stats_only = 2 [default = false];
252     repeated bytes certificates = 3;
253     optional Stats from_client = 4;
254     optional Stats from_server = 5;
255
256     optional uint32 udp_packets = 6;
257     optional uint32 tcp_packets = 7;
258     optional float udp_ping_avg = 8;
259     optional float udp_ping_var = 9;
260     optional float tcp_ping_avg = 10;
261     optional float tcp_ping_var = 11;
262
263     optional Version version = 12;
264     repeated int32 celt_versions = 13;
265     optional bytes address = 14;
266     optional uint32 bandwidth = 15;
267     optional uint32 onlinesecs = 16;
268     optional uint32 idlesecs = 17;
269     optional bool strong_certificate = 18 [default = false];
270 }
271
272 message RequestBlob {
273     repeated uint32 session_texture = 1;
274     repeated uint32 session_comment = 2;
275     repeated uint32 channel_description = 3;
276 }

```